

Platinum_Palladium_Price_Forecasting

June 22, 2020

1 Imports

```
[1]: from itertools import permutations
from functools import reduce
import itertools
import pandas as pd
import numpy as np
import pickle

# Data source
import quandl

# Facebook's Prophet
from fbprophet import Prophet

# Stats
import statsmodels.api as sm
import statsmodels.graphics.tsaplots as sgt
import statsmodels.tsa.stattools as sts
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.tsa.seasonal import STL

# Plotting
from pylab import rcParams
import matplotlib.pyplot as plt
from matplotlib.gridspec import GridSpec
plt.style.use('fivethirtyeight')

import warnings
warnings.filterwarnings("ignore")
warnings.simplefilter('once', category=UserWarning)
```

2 Functions

```
[2]: def dickey_fuller_test(df = None):
    """
    Performs the Dickey-Fuller

    Parameters
    -----
    df : dataframe
        A series of values

    Returns
    -----
    dataframe
        Returns a dataframe containing the Dickey-Fuller statistics

    """
    if df is None:
        raise RuntimeError('The dataframe and columns must both be specified')

    indices = ['Test Statistic',
               'p-value',
               'Number of lags used',
               'Number of observations',
               'Critical Value (1%)',
               'Critical Value (5%)',
               'Critical Value (10%)',
               'Maximized information criteria',
               ]

    # Perform Dickey-Fuller Test
    df_res = list(sts.adfuller(df))

    # Process output for dataframe
    stats = df_res[:4] + [x for x in df_res[4].values()] + df_res[5:]

    return pd.DataFrame({'indices': indices, 'Statistic': stats}).
    ↪set_index('indices')
```

```
[3]: def plot_acf_pacf(df, column, title, nlags=40):
    """
    Plots the observed values and their acf and pacf plots

    Parameters
    -----
    df : dataframe
        A series of values
```

```

column: string
    a string representing the column name desired from the dataframe
nlags: int
    number of lags to consider

"""
top_row = plt.subplot2grid((2,2), (0,0), rowspan=1, colspan=2)
bot_left = plt.subplot2grid((2, 2), (1, 0), rowspan=1, colspan=1)
bot_right = plt.subplot2grid((2, 2), (1, 1), rowspan=1, colspan=1)

df[column].plot(ax=top_row)
top_row.set_title(title)

sgt.plot_pacf(df[column], lags=nlags, zero=True, ax=bot_left)
bot_left.set_title('PACF')

sgt.plot_acf(df[column], lags=nlags, zero=True, ax=bot_right)
bot_right.set_title('ACF')

plt.show()

```

```

[4]: def sarima_model_selection(df, orders, method = 'lbfgs'):
    """
    Performs a grid search to find the model which minimizes the AICc score

    Parameters
    -----
    df : dataframe
        A series of values
    orders: dict
        a dictionary of lists for keys p,d,q, P, D, Q and an integer for s.
    method: string
        indicates the optimizer for the model selection

    Returns
    -----
    mod: SARIMAX model
        Returns the SARIMAX which minimizes the AICc score among all the order_
    ↪ combinations

    """

    p, d, q = orders['p'], orders['d'], orders['q']
    P, D, Q = orders['P'], orders['D'], orders['Q']
    s = orders['s']

    min_aicc = float('inf')

```

```

best_order = None
best_model = None

iterations = reduce(lambda a,b: a * b, [len(x) for x in [p,q,d,P,Q,D]])
print('There are {} combination to evaluate...'.format(iterations))

pdq = list(itertools.product(p, d, q))
seasonal_pdq = [(x[0], x[1], x[2], s) for x in list(itertools.product(P, D,
↪Q))]

for param in pdq:
    for param_seasonal in seasonal_pdq:
        try:
            mod = sm.tsa.statespace.SARIMAX(df,
                                            order=param,
                                            seasonal_order=param_seasonal,
                                            enforce_stationarity=False,
                                            enforce_invertibility=False)

            results = mod.fit(max_iter=500, disp=0)
            if results.aicc < min_aicc:
                min_aicc = results.aicc
                best_model = mod
                best_params = '{} x {}'.format(param, param_seasonal, s)
        except:
            continue
print('Best Model:{}\nLowest Score:{}'.format(best_params, min_aicc))

return mod

```

```

[5]: def plot_STL_decomposition(df, method='additive'):
    """
    Plots the observed values and their trend, seasonality, and residuals.
    ↪Additionally, outputs
    the trend strength and seasonality strength.

    Parameters
    -----
    df : dataframe
        A series of values
    method: string
        a string indicating whether to use a additive or multiplicative
    ↪decomposition method

    """
    observed_ax = plt.subplot2grid((4,1), (0,0), rowspan=1, colspan=1)
    trend_ax = plt.subplot2grid((4, 1), (1, 0), rowspan=1, colspan=1)
    season_ax = plt.subplot2grid((4, 1), (2, 0), rowspan=1, colspan=1)

```

```

resid_ax = plt.subplot2grid((4, 1), (3, 0), rowspan=1, colspan=1)

trend_strength, seasonality_strength = 0, 0

observed_ax.set_title(df.name + ' STL ' + method + ' decomposition')
observed_ax.set_ylabel('Observed')
trend_ax.set_ylabel('Trend')
season_ax.set_ylabel('Season')
resid_ax.set_ylabel('Resid')
if method == 'additive':
    stl = STL(df).fit()
    df.plot(ax=observed_ax)
    stl.trend.plot(ax=trend_ax)
    stl.seasonal.plot(ax=season_ax)
    stl.resid.plot(ax=resid_ax)
    trend_strength = max(0, 1 - np.var( stl.resid) / np.var( stl.trend +
↪stl.resid))
    seasonality_strength = max(0, 1 - np.var( stl.resid) / np.var( stl.
↪seasonal + stl.resid))
elif method == 'multiplicative':
    df = np.log(df)
    stl = STL(df).fit()
    np.exp(df).plot(ax=observed_ax)
    np.exp(stl.trend).plot(ax=trend_ax)
    np.exp(stl.seasonal).plot(ax=season_ax)
    np.exp(stl.resid).plot(ax=resid_ax)
    trend_strength = max(0, 1 - np.var(stl.resid) / np.var( stl.trend +
↪stl.resid))
    seasonality_strength = max(0, 1 - np.var( stl.resid) / np.var( stl.
↪seasonal + stl.resid))
else:
    raise RuntimeError("Method must either 'additive' or 'multiplicative'.")

print('Trend Strength: {} \nSeasonality Strength: {}'.format(trend_strength,
↪seasonality_strength))

```

3 Loading the Data

```

[6]: quandl.ApiConfig.api_key = 'k2csLJUpYjYPGKKeCSBN'
    platinum_data = quandl.get('LPPM/PLAT', column_index='1')
    palladium_prices = quandl.get('LPPM/PALL', column_index='1')

```

```

[7]: df = platinum_data.rename(columns={'USD AM': 'plat_prices'})
    df['pall_prices'] = palladium_prices['USD AM']

```

```
[8]: df.index.min(), df.index.max()
```

```
[8]: (Timestamp('1990-04-02 00:00:00'), Timestamp('2020-06-19 00:00:00'))
```

There is about 20 years worth of data

4 Pre-process Data

```
[9]: df.head()
```

```
[9]:
```

| | plat_prices | pall_prices |
|------------|-------------|-------------|
| Date | | |
| 1990-04-02 | 471.00 | 128.00 |
| 1990-04-03 | 475.80 | 128.35 |
| 1990-04-04 | 475.70 | 128.35 |
| 1990-04-05 | 481.75 | 128.40 |
| 1990-04-06 | 481.00 | 128.75 |

```
[10]: df.isna().sum()
```

```
[10]: plat_prices    0
      pall_prices    3
      dtype: int64
```

4.0.1 Fill missing data and resample into monthly frequency

```
[11]: df['pall_prices'] = df['pall_prices'].fillna(method='ffill')
      df = df[['plat_prices', 'pall_prices']].resample('MS').mean()
      df = df.apply(lambda x: round(x, 2))
```

5 Platinum

5.1 Time Plot of Platinum


```
[12]: rcParams['figure.figsize'] = 18, 8
      df.plat_prices.plot()
```

```
[12]: <matplotlib.axes._subplots.AxesSubplot at 0x7f82022d2050>
```



The time-series has a seasonally pattern, although it is hard to discern, where the prices has a series a consecutive increases followed by consecutive decreases. The trend seems to gradually until 2002 where the price spikes and experiences great flucation, likely due to the stock market crash, and then gradually decreases.

```
[13]: plot_acf_pacf(df, 'plat_prices', 'Platinum Prices')
```



While classical decomposition is widely used, it's not recommended. Instead, I will opt to use STL decomposition which is not only versatile and robust. STL has many advantages over over the classical, SEATS and X11 decomposition methods.

The variation in the seasonal pattern appears proportional to the level of the time series, therefore we will opt for a multiplicative decomposition. The caveat of STL is that it only works for additive models. However, it is possible to obtain a multiplicative decomposition by takign the logs of the data and then back-transforming the componenets. This is possible because $y_t = S_t \times T_t \times R_t \equiv$


$\log(y_t) \log(S_t) \times \log(T_t) \times \log(R_t)$

<https://otexts.com/fpp3/stl.html>

```
[14]: plot_STL_decomposition(df.plat_prices, method='multiplicative')
```

Trend Strength: 0.9932297077382147

Seasonality Strength: 0.4369190761188051



The plot indicates a trend that gradually increases until 2008 and then gradually decreases. Moreover, the plot definitely indicates seasonality

```
[15]: dickey_fuller_test(df.plat_prices)
```

| [15]: | Statistic |
|--------------------------------|-------------|
| indicies | |
| Test Statistic | -1.316759 |
| p-value | 0.621471 |
| Number of lags used | 8.000000 |
| Number of observations | 354.000000 |
| Critical Value (1%) | -3.448958 |
| Critical Value (5%) | -2.869739 |
| Critical Value (10%) | -2.571138 |
| Maximized information critiera | 3759.805843 |

Examining the Dickey-Fuller test, the p-value suggests 62% of not rejecting the null hypothesis. Additionally, the test statistic is larger than the critical values at 1%, 5%, and 10%. All this suggests we can't confirm stationarity.

ARIMA models require stationarity, so we must perform differencing

```
[16]: dickey_fuller_test(df.plat_prices.diff()[1:])
```

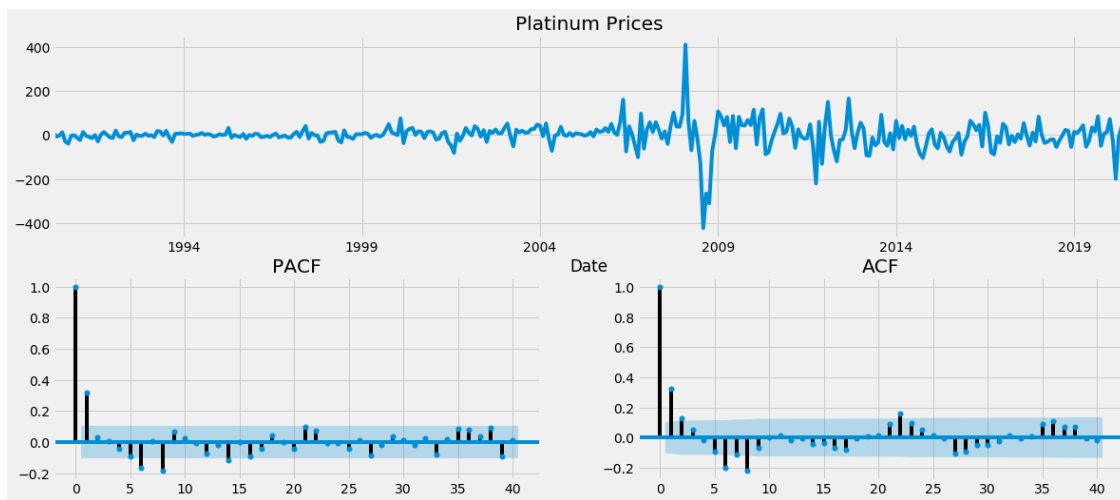


```
[16]:
```

| | Statistic |
|--------------------------------|---------------|
| indicies | |
| Test Statistic | -8.882680e+00 |
| p-value | 1.305076e-14 |
| Number of lags used | 7.000000e+00 |
| Number of observations | 3.540000e+02 |
| Critical Value (1%) | -3.448958e+00 |
| Critical Value (5%) | -2.869739e+00 |
| Critical Value (10%) | -2.571138e+00 |
| Maximized information critiera | 3.749959e+03 |

Perfect, the test statistic is much smaller than all the critical values and the p-value is also significant indicating stationarity.

```
[17]: plot_acf_pacf(df.diff()[1:], 'plat_prices', 'Platinum Prices')
```



For the AR component p , we examine the significant lags in the PACF plot. For the seasonal AR component, we look at lags of multiples of 12. We see the 24th lag is significant so we can examine seasonality when $P = 2$.

For the MA component q , we examine the significant lags in the ACF plot. For the seasonal MA component, we look at the lags of multiples of 12. There doesn't seem to be any significant lags.

```
[18]: orders = {'p': [1,2,6,8],
               'd': [1],
               'q': [1,2,3,6,8],
               'P': [0,1,2],
               'D': [1],
               'Q': [0,1],
               's': 12
            }
```

```
plat_best_model = sarima_model_selection(df.plat_prices, orders)
```

There are 120 combination to evaluate...

```
/opt/anaconda3/envs/mlenv/lib/python3.7/site-
packages/statsmodels/base/model.py:568: ConvergenceWarning: Maximum Likelihood
optimization failed to converge. Check mle_retvals
"Check mle_retvals", ConvergenceWarning)
```

Best Model:(8, 1, 1) x (2, 1, 1, 12)12

Lowest Score:3501.4581092542307

```
[19]: plat_results = plat_best_model.fit()
print(plat_results.summary().tables[1])
```

| | coef | std err | z | P> z | [0.025 | 0.975] |
|----------|-----------|---------|---------|-------|----------|----------|
| ar.L1 | -1.3628 | 0.104 | -13.094 | 0.000 | -1.567 | -1.159 |
| ar.L2 | -1.1127 | 0.183 | -6.088 | 0.000 | -1.471 | -0.754 |
| ar.L3 | -1.0501 | 0.204 | -5.142 | 0.000 | -1.450 | -0.650 |
| ar.L4 | -0.7624 | 0.227 | -3.354 | 0.001 | -1.208 | -0.317 |
| ar.L5 | -0.7762 | 0.220 | -3.536 | 0.000 | -1.206 | -0.346 |
| ar.L6 | -0.8390 | 0.181 | -4.625 | 0.000 | -1.195 | -0.483 |
| ar.L7 | -0.9734 | 0.144 | -6.754 | 0.000 | -1.256 | -0.691 |
| ar.L8 | -0.6362 | 0.084 | -7.599 | 0.000 | -0.800 | -0.472 |
| ma.L1 | 1.6726 | 0.109 | 15.317 | 0.000 | 1.459 | 1.887 |
| ma.L2 | 1.6350 | 0.205 | 7.975 | 0.000 | 1.233 | 2.037 |
| ma.L3 | 1.6614 | 0.255 | 6.511 | 0.000 | 1.161 | 2.162 |
| ma.L4 | 1.4228 | 0.282 | 5.046 | 0.000 | 0.870 | 1.975 |
| ma.L5 | 1.3915 | 0.275 | 5.061 | 0.000 | 0.853 | 1.930 |
| ma.L6 | 1.2932 | 0.233 | 5.540 | 0.000 | 0.836 | 1.751 |
| ma.L7 | 1.3362 | 0.164 | 8.144 | 0.000 | 1.015 | 1.658 |
| ma.L8 | 0.8389 | 0.082 | 10.229 | 0.000 | 0.678 | 1.000 |
| ar.S.L12 | -0.2094 | 0.101 | -2.077 | 0.038 | -0.407 | -0.012 |
| ar.S.L24 | -0.0390 | 0.094 | -0.414 | 0.679 | -0.224 | 0.146 |
| ma.S.L12 | -0.8854 | 0.073 | -12.075 | 0.000 | -1.029 | -0.742 |
| sigma2 | 4087.9684 | 419.815 | 9.738 | 0.000 | 3265.146 | 4910.791 |

Many of the p-values are significant indicating the lags are significantly different from 0.

5.1.1 Ljung-Box test

In addition to looking at the ACF plot, we can do a more formal test for autocorrelation by considering a whole set residuals instead of treating each one separately. Therefore, we test whether the first autocorrelations are significantly different from what would be expected from a white noise process.

It is suggested to use lags = 10 for non-seasonal data and lags = 2m for seasonal data where m is the period of seasonality. However, if

However, the test is not good when the number of lags is too high so if the the number of lags is larger than $T/5$ where T is the number of observations, then use lags = $T/5$.


source: <https://otexts.com/fpp3/diagnostics.html>

```
[20]: sm.stats.acorr_ljungbox(plat_results.resid, lags=[24], return_df=True)
```

```
[20]:      lb_stat  lb_pvalue
      24  57.705438   0.000134
```


Our p-value indicates the first autocorrelations are significantly different from what would be expected from a white noise process.

```
[21]: plat_results.plot_diagnostics(figsize=(16, 8))
      plt.show()
```



```
[44]: pred = plat_results.get_prediction(start=pd.to_datetime('2019-06-01'),
    ↪dynamic=False)
pred_ci = pred.conf_int()
ax = df.plat_prices['2016:'].plot(label='Observed')
pred.predicted_mean.plot(ax=ax, label='One-step ahead Forecast', alpha=.7,
    ↪figsize=(14, 7))
ax.fill_between(pred_ci.index,
                pred_ci.iloc[:, 0],
                pred_ci.iloc[:, 1], color='k', alpha=.2)
ax.set_xlabel('Date')
ax.set_ylabel('Platinum Sales')
```

```
plt.legend()
plt.show()
```




The forecasts seem to align well with the observed values.

```
[23]: y_forecasted = pred.predicted_mean
y_truth = df.plat_prices['2008-01-01':]
mse = ((y_forecasted - y_truth) ** 2).mean()
print('The Mean Squared Error of our forecasts is {}'.format(round(mse, 2)))
print('The Root Mean Squared Error of our forecasts is {}'.format(round(np.
    ↳sqrt(mse), 2)))
```

The Mean Squared Error of our forecasts is 3536.01

The Root Mean Squared Error of our forecasts is 59.46

```
[43]: pred_uc = plat_results.get_forecast(steps=36)
pred_ci = pred_uc.conf_int()
ax = df.plat_prices['2015-07-01:'].plot(label='observed', figsize=(14, 7))
pred_uc.predicted_mean.plot(ax=ax, label='Forecast')
ax.fill_between(pred_ci.index,
                pred_ci.iloc[:, 0],
                pred_ci.iloc[:, 1], color='k', alpha=.25)
ax.set_xlabel('Date')
ax.set_ylabel('Platinum Prices')
plt.legend()
plt.show()
```



6 Palladium

6.1 Time Plot of Palladium

```
[25]: rcParams['figure.figsize'] = 18, 8
      df.pall_prices.plot()
```

```
[25]: <matplotlib.axes._subplots.AxesSubplot at 0x7f81d81cf6d0>
```



The time-series has a seasonally pattern with occasional spikes such as in 2001-2002. The trend increases through the entire timeseries with notable exponential growth after 2016.


```
[26]: dickey_fuller_test(df.pall_prices)
```

```
[26]:
```

| | Statistic |
|--------------------------------|-------------|
| indicies | |
| Test Statistic | 0.812363 |
| p-value | 0.991833 |
| Number of lags used | 17.000000 |
| Number of observations | 345.000000 |
| Critical Value (1%) | -3.449447 |
| Critical Value (5%) | -2.869954 |
| Critical Value (10%) | -2.571253 |
| Maximized information critiera | 3691.410626 |

The p-value suggests 99.2% of not rejecting the null hypothesis. Additionally, the test statistic is larger than the critical values at 1%, 5%, and 10%. All this suggests we can't confirm stationarity


```
[27]: plot_acf_pacf(df, 'pall_prices', 'Palladium Prices')
```



As above, we will use multiplicative STL decomposition

```
[28]: stl = plot_STL_decomposition(df.pall_prices, method='multiplicative')
```

Trend Strength: 0.9940630024022675
Seasonality Strength: 0.4592015311752917



Let's examine the stationary after differencing


```
[29]: dickey_fuller_test(df.pall_prices.diff()[1:])
```

```
[29]:
```

| | Statistic |
|--------------------------------|-------------|
| indices | |
| Test Statistic | -3.484995 |
| p-value | 0.008381 |
| Number of lags used | 12.000000 |
| Number of observations | 349.000000 |
| Critical Value (1%) | -3.449227 |
| Critical Value (5%) | -2.869857 |
| Critical Value (10%) | -2.571201 |
| Maximized information critiera | 3679.906205 |

Perfect, the test statistic is much smaller than all the critical values and the p-value is also significant indicating stationarity.

```
[30]: plot_acf_pacf(df.diff()[1:], 'pall_prices', 'Palladium Prices')
```



```
[31]: orders = {'p': [1,2, 9], #8
               'd': [1],
               'q': [1,2, 9],
               'P': [0,1],
               'D': [1],
               'Q': [0,1],
               's': 12
        }

pall_best_model = sarima_model_selection(df.pall_prices, orders)
```

There are 36 combination to evaluate...

```
/opt/anaconda3/envs/mlenv/lib/python3.7/site-
packages/statsmodels/base/model.py:568: ConvergenceWarning: Maximum Likelihood
optimization failed to converge. Check mle_retvals
  "Check mle_retvals", ConvergenceWarning)
```

Best Model:(2, 1, 9) x (1, 1, 1, 12)12

Lowest Score:3546.042914254437

Best Model:(1, 1, 9) x (0, 1, 1, 12)12 Lowest Score:3549.1285862915197

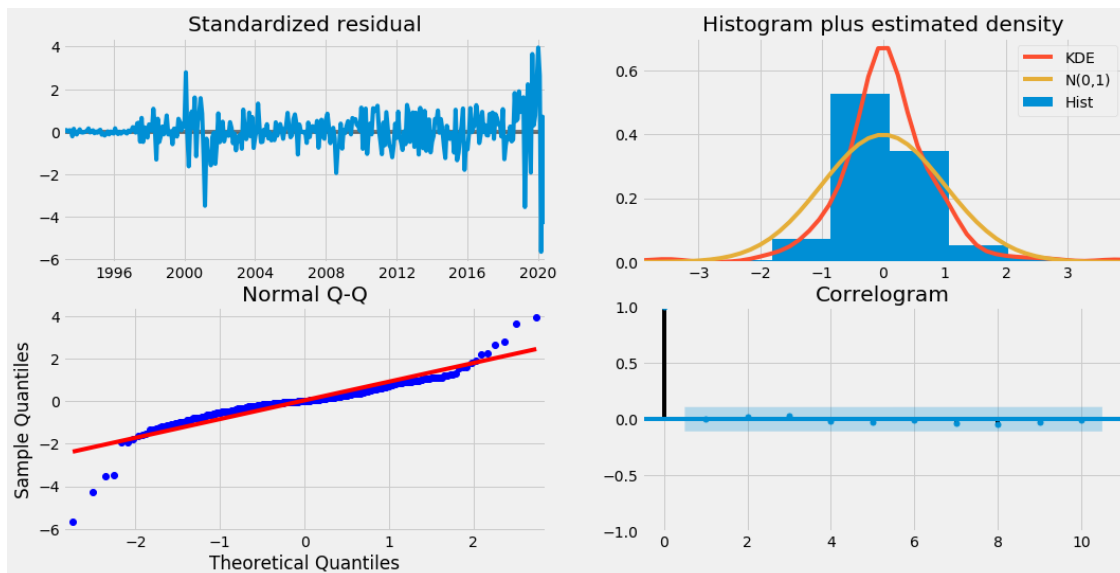
```
[32]: pall_results = pall_best_model.fit()
       print(pall_results.summary().tables[1])
```

```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          0.1705        0.659        0.259      0.796       -1.122        1.463
ar.L2         -0.3624        0.565       -0.641      0.522       -1.471        0.746
ar.L3         -0.7616        0.600       -1.269      0.204       -1.938        0.415
```


| | | | | | | |
|----------|-----------|---------|--------|-------|----------|----------|
| ar.L4 | 0.1224 | 0.905 | 0.135 | 0.892 | -1.652 | 1.897 |
| ar.L5 | -0.1744 | 0.788 | -0.221 | 0.825 | -1.719 | 1.370 |
| ar.L6 | 0.2941 | 0.573 | 0.514 | 0.607 | -0.828 | 1.416 |
| ar.L7 | 0.0770 | 0.466 | 0.165 | 0.869 | -0.837 | 0.991 |
| ar.L8 | -0.0114 | 0.384 | -0.030 | 0.976 | -0.764 | 0.742 |
| ar.L9 | 0.5283 | 0.298 | 1.776 | 0.076 | -0.055 | 1.111 |
| ma.L1 | 0.0686 | 0.664 | 0.103 | 0.918 | -1.233 | 1.371 |
| ma.L2 | 0.3112 | 0.672 | 0.463 | 0.643 | -1.006 | 1.629 |
| ma.L3 | 0.7651 | 0.680 | 1.125 | 0.260 | -0.568 | 2.098 |
| ma.L4 | 0.1757 | 0.990 | 0.177 | 0.859 | -1.764 | 2.116 |
| ma.L5 | 0.2270 | 1.028 | 0.221 | 0.825 | -1.788 | 2.242 |
| ma.L6 | -0.4673 | 0.860 | -0.543 | 0.587 | -2.153 | 1.219 |
| ma.L7 | 0.2207 | 0.551 | 0.400 | 0.689 | -0.860 | 1.301 |
| ma.L8 | -0.1454 | 0.606 | -0.240 | 0.810 | -1.332 | 1.041 |
| ma.L9 | -0.5643 | 0.447 | -1.262 | 0.207 | -1.441 | 0.312 |
| ar.S.L12 | 0.0133 | 0.158 | 0.084 | 0.933 | -0.296 | 0.323 |
| ma.S.L12 | -0.8855 | 0.145 | -6.126 | 0.000 | -1.169 | -0.602 |
| sigma2 | 3264.7086 | 248.572 | 13.134 | 0.000 | 2777.517 | 3751.900 |

=====

```
[33]: pall_results.plot_diagnostics(figsize=(16, 8))
plt.show()
```




```
[34]: pred = pall_results.get_prediction(start=pd.to_datetime('2019-06-01'),
    ↪dynamic=False)
pred_ci = pred.conf_int()
ax = df.pall_prices['2016:'].plot(label='Observed')
```

```

pred.predicted_mean.plot(ax=ax, label='One-step ahead Forecast', alpha=.7,
    figsize=(14, 7))
ax.fill_between(pred_ci.index,
                pred_ci.iloc[:, 0],
                pred_ci.iloc[:, 1], color='k', alpha=.2)
ax.set_xlabel('Date')
ax.set_ylabel('Palladium Sales')
plt.legend()
plt.show()

```



The forecasts seem to align well with the observed values

```

[35]: y_forecasted = pred.predicted_mean
y_truth = df.pall_prices['2008-01-01':]
mse = ((y_forecasted - y_truth) ** 2).mean()
print('The Mean Squared Error of our forecasts is {}'.format(round(mse, 2)))
print('The Root Mean Squared Error of our forecasts is {}'.format(round(np.
    sqrt(mse), 2)))

```

The Mean Squared Error of our forecasts is 25437.19


The Root Mean Squared Error of our forecasts is 159.49

```

[36]: pred_uc = pall_results.get_forecast(steps=36)
pred_ci = pred_uc.conf_int()
ax = df.pall_prices['2015-07-01':].plot(label='observed', figsize=(14, 7))
pred_uc.predicted_mean.plot(ax=ax, label='Forecast')
ax.fill_between(pred_ci.index,
                pred_ci.iloc[:, 0],
                pred_ci.iloc[:, 1], color='k', alpha=.25)

```

```
ax.set_xlabel('Date')
ax.set_ylabel('Palladium Sales')
plt.legend()
plt.show()
```




7 Prophet


```
[37]: fbp_plat_model = Prophet()
fbp_plat_model.fit(pd.DataFrame({'ds': df.index, 'y': df.plat_prices}))
fbp_plat_future = fbp_plat_model.make_future_dataframe(periods=24, freq='MS')
fbp_plat_forecast = fbp_plat_model.predict(fbp_plat_future)
fbp_plat_model.plot(fbp_plat_forecast, xlabel = 'Date', ylabel = 'Prices')
plt.show()
```

INFO:fbprophet:Disabling weekly seasonality. Run prophet with weekly_seasonality=True to override this.

INFO:fbprophet:Disabling daily seasonality. Run prophet with daily_seasonality=True to override this.




```
[38]: fbp_plat_model.plot_components(fbp_plat_forecast);
```




```
[39]: fbp_pall_model = Prophet()
fbp_pall_model.fit(pd.DataFrame({'ds': df.index, 'y': df.pall_prices}))
fbp_pall_future = fbp_pall_model.make_future_dataframe(periods=24, freq='MS')
fbp_pall_forecast = fbp_pall_model.predict(fbp_plat_future)
fbp_pall_model.plot(fbp_pall_forecast, xlabel = 'Date', ylabel = 'Prices')
plt.show()
```

INFO:fbprophet:Disabling weekly seasonality. Run prophet with weekly_seasonality=True to override this.


INFO:fbprophet:Disabling daily seasonality. Run prophet with daily_seasonality=True to override this.




```
[40]: fbp_pall_model.plot_components(fbp_pall_forecast);
```



```
[41]: plt.plot(fbp_plat_forecast['ds'], fbp_plat_forecast['trend'], label='Platinum')
plt.plot(fbp_pall_forecast['ds'], fbp_pall_forecast['trend'], label='Palladium')
plt.xlabel('Date')
plt.ylabel('Prices')
plt.title('Platinum vs. Palladium Price Trend');
plt.legend()
plt.show()
```



```
[42]: plt.plot(fbp_plat_forecast['ds'], fbp_plat_forecast['yhat'], label='Platinum')
plt.plot(fbp_pall_forecast['ds'], fbp_pall_forecast['yhat'], label='Paladium')
plt.xlabel('Date')
plt.ylabel('Prices')
plt.title('Platinum vs. Palladium Price Estimates');
plt.legend()
plt.show()
```



7.1 Pickle Forecasting Models

```
[43]: pickle.dump( plat_results, open( "platinum_forecasting_model.p", "wb" ) )
pickle.dump( pall_results, open( "palladium_forecasting_model.p", "wb" ) )
pickle.dump( fbp_plat_model, open( "fbprophet_platinum_forecasting_model.p", "wb" ) )
pickle.dump( fbp_pall_model, open( "fbprophet_palladium_forecasting_model.p", "wb" ) )
```